
§2.1 Solving Linear Systems

- Augmented coefficient matrix:

$$A_{11}x_1 + A_{12}x_2 + \dots + A_{1n}x_n = b_1$$

$$A_{21}x_1 + A_{22}x_2 + \dots + A_{2n}x_n = b_2$$

....

$$A_{n1}x_1 + A_{n2}x_2 + \dots + A_{nn}x_n = b_n$$

Uniqueness of Solution

- Unique if \mathbf{A} is non-singular or $\det(\mathbf{A}) = |\mathbf{A}| \neq 0$; equivalently, the rows of \mathbf{A} and columns of \mathbf{A} must be *linearly* independent.
- If \mathbf{A} is singular, then $\mathbf{Ax}=\mathbf{b}$ has ∞ or no solutions.

Ill-conditioning

- Case when coefficient matrix A is *nearly* singular (or when $|A|$ is very small).
- Can estimate $|A| \ll \|A\|$, where $\|A\|$ is a matrix norm:

Ill-conditioning

- The condition number of a \mathbf{A} is defined by $\text{cond}(\mathbf{A}) = ||\mathbf{A}||_{\times} ||\mathbf{A}^{-1}||$.
- If $\text{cond}(\mathbf{A}) \approx 1$, then we say \mathbf{A} is well-conditioned; $\text{cond}(\mathbf{A})$ is not unique – depends on the norm; $\text{cond}(\mathbf{A})$ is expensive to compute so we typically compare $||\mathbf{A}||$ to magnitudes of matrix \mathbf{A} elements.

Ill-conditioning

- Suppose we modify the second equation of the previous example to be $2x + 1.002y = 0$. What does the solution to the 2 by 2 linear system become?
- So, we see that a _____% change in A yields a _____% change in the solution.

Methods of Solution

- **Direct methods** – based on *elementary operations* that do not augment the solution:
 1. Exchange 2 rows (changes sign of $|A|$);
 2. Multiply equation by nonzero constant α ($|A|$ becomes $\alpha |A|$);
 3. Multiply equation by nonzero constant α and subtract it from another equation ($|A|$ unchanged).
- **Iterative methods** (or indirect methods) – typically used for large and sparse A (more zeros than nonzeros); guess solution and improve every iteration.

Direct Methods

<u>Method</u>	<u>Initial Form</u>	<u>Final Form</u>
Gauss Elimination	$Ax=b$	$Ux=c$
LU Decomposition	$Ax=b$	$LUx=b$
Gauss-Jordan Elim.	$Ax=b$	$Ix=c$

Sample U,L for 3-by-3 matrices:

LU Decomposition

- Transform $Ax=b$ given $A=LU$:
 $(LU)x=b \rightarrow L(Ux)=b$; let $Ux=y$ and solve $Ly=b$ for y via *forward substitution*, then solve $Ux=y$ for the solution vector x .

LU Decomposition

- Now that we have $y=(y_1 \ y_2 \ y_3)^T$ use *back-substitution* to get $x=(x_1 \ x_2 \ x_3)^T$:

LU Decomposition Example

$$A = \begin{pmatrix} 8 & -6 & 2 \\ -4 & 11 & -7 \\ 4 & -7 & 6 \end{pmatrix}, \quad b = \begin{pmatrix} 28 \\ -40 \\ 33 \end{pmatrix}$$

Gaussian Elimination

- Use elementary row operations that preserve the solution and produce an *upper-triangular* augmented coefficient matrix.

Example:
$$\begin{array}{rcl} 4x_1 - 2x_2 + x_3 & = & 11 \\ -2x_1 + 4x_2 - 2x_3 & = & -16 \\ x_1 - 2x_2 + 4x_3 & = & 17 \end{array} \rightarrow \left(\begin{array}{ccc|c} 4 & -2 & 1 & 11 \\ -2 & 4 & -2 & -16 \\ 1 & -2 & 4 & 17 \end{array} \right) \begin{array}{l} \mathbf{a} \\ \mathbf{b} \\ \mathbf{c} \end{array}$$

Gaussian Elimination

- Now element in (2,2) position is the *pivot*:

$$Eq(c) \leftarrow Eq(c) - (-0.5) \times Eq(b) \rightarrow \left(\begin{array}{ccc|c} 4 & -2 & 1 & 11 \\ 0 & 3 & -1.5 & -10.5 \\ 0 & 0 & 3 & 9 \end{array} \right)$$

$$(x_1, x_2, x_3)^T = (\underline{\quad}, \underline{\quad}, \underline{\quad})^T$$

Gaussian Elimination

- Python code for elimination phase (getting the upper-triangular augmented coefficient matrix):

```
for k in range(0,n-1):  
    for i in range(k+1,n):  
        if a[i,k] != 0.0  
            lam=a[i,k]/a[k,k]      #pivot  
            a[i,k+1:n]=a[i,k+1:n]-lam*a[k,k+1:n]  
            b[i]=b[i]-lam*b[k]
```

- What happens to A_{ij} for $i > j$?

Gaussian Elimination

- See `gaussElim.py` on page 41 of textbook.
- Can handle multiple right-hand-sides ($\mathbf{AX}=\mathbf{B}$) with minor changes to `gaussElim.py` but method would be inefficient; *back-substitution* has to be repeated for each final column of \mathbf{B} .

§2.3 LU Decomposition

- Can express any n -by- n matrix A as $A=LU$.

Different forms:

Doolittle $L_{ii}=1, i=1,2,\dots,n$

Crout $U_{ii}=1, i=1,2,\dots,n$

Choleski $L=U^T$ (for symmetric matrices A)

How do we solve $Ax=b$ using $A=LU$?

Doolittle's Decomposition

- Most common LU decomposition; matrix (or array \mathbf{A}) can be overwritten:
Implicit that the diagonal of \mathbf{L} is all ones; element in \mathbf{L}_{ij} is the pivot equation multiplier λ needed to zero the current \mathbf{A}_{ij} element; LUdecomp.py on p. 47 in textbook.

$$\begin{pmatrix} U_{11} & U_{12} & U_{13} \\ L_{21} & U_{22} & U_{23} \\ L_{31} & L_{32} & U_{33} \end{pmatrix}$$

Choleski Decomposition

- $\mathbf{A} = \mathbf{L}\mathbf{L}^T$
- \mathbf{A} must be symmetric (i.e., $\mathbf{A} = \mathbf{A}^T$)
- \mathbf{A} must be positive definite ($\vec{x}^T \mathbf{A} \vec{x} > 0 \ \forall \vec{x} \neq \vec{0}$)

$$\begin{pmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix} = \begin{pmatrix} L_{11} & 0 & 0 \\ L_{21} & L_{22} & 0 \\ L_{31} & L_{32} & L_{33} \end{pmatrix} \begin{pmatrix} L_{11} & L_{21} & L_{31} \\ 0 & L_{22} & L_{32} \\ 0 & 0 & L_{33} \end{pmatrix}$$

Choleski Decomposition

- Match A to the product LL^T :

$$\begin{pmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix} = \begin{pmatrix} L_{11}^2 & L_{11}L_{21} & L_{11}L_{31} \\ L_{11}L_{21} & L_{21}^2 + L_{22}^2 & L_{21}L_{31} + L_{22}L_{32} \\ L_{11}L_{31} & L_{21}L_{31} + L_{22}L_{32} & L_{31}^2 + L_{32}^2 + L_{33}^2 \end{pmatrix}$$

Choleski Decomposition

- Generalization for n -by- n matrices:

$$L_{jj} = \sqrt{A_{jj} - \sum_{k=1}^{j-1} L_{jk}^2}, \quad j = 2, 3, \dots, n$$

$$L_{ij} = (A_{ij} - \sum_{k=1}^{j-1} L_{ik} L_{jk}) / L_{jj}, \quad j = 2, 3, \dots, n; i = j+1, j+2, \dots, n$$

- See `choleski.py` on p. 50 of textbook.
- Applications: Example 2.8 (p. 54), Problem Set 2.1 (#16, p. 57)

§2.4 Symmetric & Banded Matrices

- Sparse matrix – most elements are zero.
- Banded matrix – nonzeros are clustered near the main diagonal; tridiagonal matrix **A** has a *bandwidth* of 3:

$$A = \begin{pmatrix} x & x & & & \\ x & x & x & & \\ & x & x & x & \\ & & x & x & x \\ & & & x & x \end{pmatrix}$$

Banded Matrices

- For a banded matrix \mathbf{A} with $\mathbf{A}=\mathbf{LU}$, both \mathbf{L} and \mathbf{U} retain the banded *structure* of \mathbf{A} :

$$\mathbf{A} = \begin{pmatrix} x & x & & & \\ x & x & x & & \\ & x & x & x & \\ & & x & x & x \\ & & & x & x \end{pmatrix}, \quad \mathbf{L} = \begin{pmatrix} x & & & & \\ x & x & & & \\ & x & x & & \\ & & x & x & \\ & & & x & x \end{pmatrix}, \quad \mathbf{U} = \begin{pmatrix} x & x & & & \\ & x & x & & \\ & & x & x & \\ & & & x & x \\ & & & & x \end{pmatrix}$$

Banded Matrices

- How can we exploit symmetry & banded structure in solving $Ax=b$?
- For *tridiagonal* coefficient matrices, we only need to eliminate one element below each pivot element, i.e., only one elementary operation is needed per pivot row.
- See `LUdecomp3.py` on pp. 61-62 of the textbook.

Banded Matrices

- In `LUdecomp3.py` notice that `b` is overwritten with the solution vector in `LUsolve3(c,d,e,b)`.
- Application: Example 2.11 on p.68 of the textbook.

Symmetric Coefficient Matrices

- It is common to have symmetric and / or banded matrices in engineering applications; $\mathbf{A}=\mathbf{A}^T$ but not necessarily positive definite in this case.

If $\mathbf{A}=\mathbf{A}^T$, then $\mathbf{A}=\mathbf{LU}=\mathbf{LDL}^T$, where \mathbf{D} is a diagonal matrix.

- We can use `LUdecomp.py` and recover \mathbf{L} and \mathbf{D} from the \mathbf{U} factor.

Symmetric Coefficient Matrices

- Matching elements of \mathbf{U} we can recover \mathbf{L} and \mathbf{D} :

$$\begin{pmatrix} D_1 & & & & \\ & D_2 & & & \\ & & D_3 & & \\ & & & \dots & \\ & & & & D_n \end{pmatrix} \begin{pmatrix} 1 & L_{21} & L_{31} & \dots & L_{n1} \\ & 1 & L_{32} & \dots & L_{n2} \\ & & 1 & \dots & L_{n3} \\ & & & \dots & \dots \\ & & & & 1 \end{pmatrix} = \begin{pmatrix} D_1 & D_1 L_{21} & D_1 L_{31} & \dots & D_1 L_{n1} \\ & D_2 & D_2 L_{32} & \dots & D_2 L_{n2} \\ & & D_3 & \dots & D_3 L_{n3} \\ & & & \dots & \dots \\ & & & & D_n \end{pmatrix}$$

- Application: Problem Set 2.2 (#3, p. 78): given 5-by-5 symmetric tridiagonal matrix \mathbf{A} , determine \mathbf{L} and \mathbf{D} so that $\mathbf{A} = \mathbf{L}\mathbf{D}\mathbf{L}^T$.

§2.5 Pivoting

- Need to reorder equations during the elimination phase in order to avoid very *small* multipliers (pivot elements).
- Suppose the augmented matrix is
$$\left(\begin{array}{ccc|c} \varepsilon & -1 & 1 & 0 \\ -1 & 2 & -1 & 0 \\ 2 & -1 & 0 & 1 \end{array} \right)$$

Pivoting

- We don't always have to pick A_{kk} for the pivot element when zeroing out elements in column k (below the diagonal).

- Instead, we find A_{pk} of the largest *relative* size (r_{pk}):

$$r_{pk} = \max_j (r_{jk}), j \geq k,$$

$$\text{where } r_{ij} = \frac{|A_{ij}|}{s_i}, s_i = \max_j |A_{ij}|, i = 1, 2, \dots, n.$$

Pivoting

- If $p \neq k$, we interchange rows k and p and proceed with the elimination.
- See `gaussPivot.py` on pp. 72-73; uses the `swap.py` function on p. 67.
- See also the `LUpivot.py` module on pp. 74-75; incorporates pivoting into Doolittle's decomposition and row permutations are stored in the `seq` vector.
- Sample application: Problem Set 2.2 (#12) on pp. 79-80.

Pivoting (Caveats)

- Destroys symmetry and bandedness which do arise in many engineering applications.
- Pivoting is unnecessary when the matrix is *diagonally dominant*:
$$\begin{pmatrix} 4 & -2 & 1 \\ -2 & 4 & -1 \\ 1 & -1 & 3 \end{pmatrix}$$
- Computation time is increased with pivoting.

§2.6 Matrix Inversion

- Can compute A^{-1} for an n -by- n matrix A by solving $AX=I_n$.
- If A is banded, A^{-1} will be dense (lose nonzero structure); if A is triangular, A^{-1} will also be triangular.
- Cost of inverting A to solve $Ax=b$ (i.e., $x=A^{-1}b$) is much higher than using LU decomposition.

§2.7 Iterative Methods

- Sometimes they are called *indirect* methods.
- Take an initial guess at the solution x (for $Ax=b$) and repeatedly improve x until change is negligible.
- Advantages: store only nonzeros of A and methods are self-correcting (round-off errors are corrected in subsequent iterations).

Iterative Methods

- Drawback: do not always converge to the exact solution; if the matrix A is *diagonally dominant*, convergence is **guaranteed**.
- Initial guess affects the number of iterations (not really if method converges).
- First method we consider is Gauss-Seidel.

Gauss-Seidel

- Write $Ax=b$ as $\sum_{j=1}^n A_{ij}x_j = b_i, i=1,2,\dots,n.$
- Extract term for x_i to get $A_{ii}x_i + \sum_{\substack{j=1 \\ j \neq i}}^n A_{ij}x_j = b_i, i=1,2,\dots,n.$
- Solving for x_i yields:
$$x_i = \frac{1}{A_{ii}} \left(b_i - \sum_{\substack{j=1 \\ j \neq i}}^n A_{ij}x_j \right), i=1,2,\dots,n.$$

Gauss-Seidel

- Algorithm:
$$x_i^{(k+1)} \leftarrow \frac{1}{A_{ii}} \left(b_i - \sum_{\substack{j=1 \\ j \neq i}}^n A_{ij} x_j^{(k)} \right), i = 1, 2, \dots, n.$$
- Start with an initial guess, $x^{(0)}$ and generate successive iterates $x^{(1)}, x^{(2)}, \dots$, etc. until the difference between $x^{(k+1)}$ and $x^{(k)}$ is sufficiently small; how can convergence be improved?

Gauss-Seidel

- Algorithm with relaxation:

$$x_i^{(k+1)} \leftarrow \frac{\omega}{A_{ii}} \left(b_i - \sum_{\substack{j=1 \\ j \neq i}}^n A_{ij} x_j^{(k)} \right) + (1 - \omega) x_i^{(k)}, i = 1, 2, \dots, n.$$

- Take weighted average of previous iteration values; $\omega < 1$ (under-relaxation), $\omega > 1$ (over-relaxation). What is optimal ω ?

Gauss-Seidel

- Can be shown that $\omega_{\text{opt}} = \frac{2}{1 + \sqrt{1 - \left(\frac{\Delta x^{(k+p)}}{\Delta x^{(k)}} \right)^{1/p}}}$,

where p is a positive integer and $\Delta x^{(k)} = |x^{(k-1)} - x^{(k)}|$.

- Strategy: (1) perform k iterations with $\omega=1$, after k^{th} iteration record $\Delta x^{(k)}$; (2) perform p more iterations and record $\Delta x^{(k+p)}$; and (3) perform subsequent iterations with $\omega = \omega_{\text{opt}}$; see `gaussSeidel.py` on p.89 and **Example 2.17** on pp. 95-96.

Conjugate Gradient Method

- Consider the function $f(x) = \frac{1}{2} x^T \mathbf{A}x - b^T x$, where \mathbf{A} is a symmetric and positive definite (SPD) matrix and x, b are vectors in R^n .
- The minimum for $f(x)$ occurs when the gradient $\nabla f = \mathbf{A}x - b = 0$ or when $\mathbf{A}x = b$.
- Desire an iteration of the form $x_{k+1} = x_k + \alpha_k s_k$; where s_k is a *search* direction and α_k is the *step length*.

Conjugate Gradient Method


- To minimize $f(x_{k+1})$ we choose α_k so that $\mathbf{A}(x_k + \alpha_k s_k) = b$, why?
- Let $r_k = b - \mathbf{A}x_k$ (k^{th} residual? vector) so that $\mathbf{A}(x_k + \alpha_k s_k) = b$ after multiplying through by \mathbf{A} yields $\alpha_k \mathbf{A} s_k = b - \mathbf{A}x_k = r_k$.
- Using this last equality and solving for α_k yields $\alpha_k = s_k^T r_k / s_k^T \mathbf{A} s_k$ (a way to compute the step length).
- How do we determine the search direction s_k ?

Conjugate Gradient Method

- Could choose $s_k = -\nabla f = r_k$ (i.e., direction of largest negative change in $f(x)$); this is called the method of *steepest descent*.
- How would it converge?
- Alternative approach: conjugate gradient with $s_{k+1} = r_{k+1} + \beta_k s_k$, where the constant β_k is chosen so that two successive s_k 's are *conjugate*. This means $s_{k+1}^T A s_k = 0$.

Conjugate Gradient Method

- CG preserves minimizations from previous iterations (no backtracking); not the case for steepest descent.
- Compare convergence paths:



Assuming **exact** arithmetic, CG would converge in at most n steps for an n -by- n linear system.

Conjugate Gradient Method

- Since $s_{k+1} = r_{k+1} + \beta_k s_k$ and we require $s_{k+1}^T A s_k = 0$, then we must have $(r_{k+1} + \beta_k s_k)^T A s_k = 0$.
- Solving for β_k we obtain... $\beta_k = -r_{k+1}^T A s_k / s_k^T A s_k$ (and we can now advance the search direction).
- Initialization of CG algorithm:
 - 1) Choose x_0
 - 2) Compute $r_0 = b - A x_0$
 - 3) Choose $s_0 = r_0$ (i.e., start with direction of steepest descent)

Conjugate Gradient Method

- Main loop of CG algorithm:
 - 4) For $k = 0, 1, 2, \dots$
 - $\alpha_k = s_k^T r_k / s_k^T A s_k$
 - $x_{k+1} = x_k + \alpha_k s_k$
 - $r_{k+1} = b - A x_{k+1}$
 - If $|r_{k+1}| \leq \varepsilon$, exit loop. (ε is the error tolerance.)
 - $\beta_k = -r_{k+1}^T A s_k / s_k^T A s_k$
 - $s_{k+1} = r_{k+1} + \beta_k s_k$
- Residual vectors r_1, r_2, r_3, \dots are *mutually orthogonal* (i.e., $r_i^T r_j = 0$ for $i \neq j$).

Conjugate Gradient Method

- Since the n residual vectors $\{r_1, r_2, r_3, \dots, r_n\}$ are *mutually orthogonal*, we know that $r_{n+1} = 0$. Why?
- This means that CG should theoretically converge in n iterations of the loop; typically obtain convergence in much less than n iterations.
- See `conjGrad.py` on pp.91-92 of textbook and review Example 2.18 on p.97.